

COMP9313: Big Data Management

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



Spark

Lecturer: Xin Cao

Course web site: <http://www.cse.unsw.edu.au/~cs9313/>

Assignment Project Exam Help

Chapter 5: Graph Data Processing

<https://powcoder.com>

In MapReduce

Add WeChat powcoder

What's a Graph?

- $G = (V, E)$, where
 - V represents the set of vertices (nodes)
 - E represents the set of edges (links)
 - Both vertices and edges may contain additional information
- Different types of graphs
 - Directed vs. undirected edges
 - Presence or absence of cycles
- Graphs are everywhere:
 - Hyperlink structure of the Web
 - Physical structure of computers on the Internet
 - Interstate highway system
 - Social networks

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

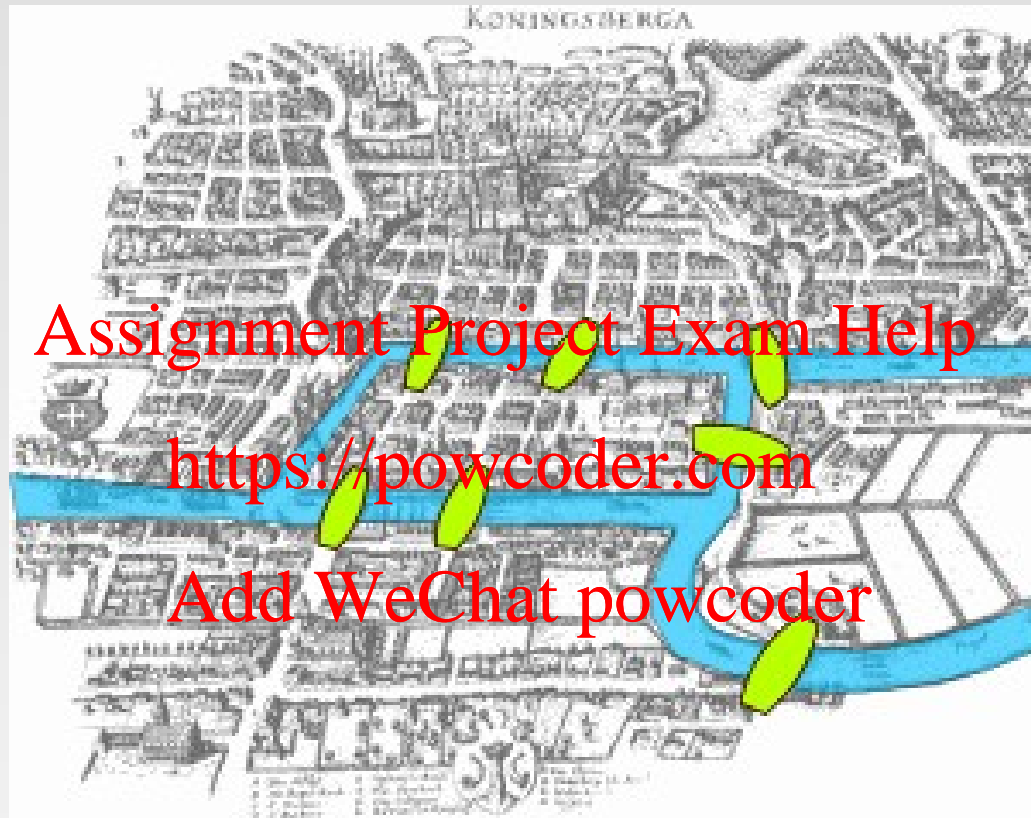
Graph Data: Social Networks



Facebook social graph

4-degrees of separation [Backstrom-Boldi-Rosa-Ugander-Vigna, 2011]

Graph Data: Technological Networks



Assignment Project Exam Help


<https://powcoder.com>

Add WeChat powcoder

Seven Bridges of Königsberg

[Euler, 1735]

Return to the starting point by traveling each link of the graph once and only once.



Some Graph Problems

- Finding shortest paths
 - Routing Internet traffic and UPS trucks
- Finding minimum spanning trees
 - Telco laying down fiber
- Finding Max Flow
 - Airline scheduling
- Identify “special” nodes and communities
 - Breaking up terrorist cells, spread of avian flu
- Bipartite matching
 - Monster.com, Match.com
- And of course... PageRank

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Graph Analytics

- General Graph
 - Count the number of nodes whose degree is equal to 5
 - Find the diameter of the graphs
- Web Graph
 - Rank each web page in the web graph or each user in the twitter graph using PageRank, or other centrality measure
- Transportation Network
 - Return the shortest or cheapest flight/road from one city to another
- Social Network
 - Detect a group of users who have similar interests
- Financial Network
 - Find the path connecting two suspicious transactions;
-

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Graphs and MapReduce

- Graph algorithms typically involve:
 - Performing computations at each node: based on node features, edge features, and local link structure
 - Propagating computations: “traversing” the graph
- Key questions:
 - How do you represent graph data in MapReduce?
 - How do you traverse a graph in MapReduce?

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Representing Graphs


- Adjacency Matrices: Represent a graph as an $n \times n$ square matrix M
 - $n = |V|$
 - $M_{ij} = 1$ means a link from node i to j

Assignment Project Exam Help

	1	2	3	4
1	0	1	0	1
2	1	0	1	1
3	1	0	0	0
4	1	0	1	0

<https://powcoder.com>

Add WeChat powcoder



Adjacency Matrices: Critique

- Advantages:
 - Amenable to mathematical manipulation
 - Iteration over rows and columns corresponds to computations on outlinks and inlinks
- Disadvantages:
 - Lots of zeros for sparse matrices
 - Lots of wasted space

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Representing Graphs

- Adjacency Lists: Take adjacency matrices... and throw away all the zeros

Assignment Project Exam Help

	1	2	3	4
1	0	1	0	1
2	1	0	1	1
3	1	0	0	0
4	1	0	1	0

<https://powcoder.com>

Add WeChat powcoder



1: 2, 4

2: 1, 3, 4

3: 1

4: 1, 3

Adjacency Lists: Critique

- Advantages:

- Much more compact representation
- Easy to compute over outlinks

- Disadvantages:

- Much more difficult to compute over inlinks

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder


Assignment Project Exam Help

<https://powcoder.com>
Single-Source Shortest Path

Add WeChat powcoder

Single-Source Shortest Path (SSSP)

- **Problem:** find shortest path from a source node to one or more target nodes
 - Shortest might also mean lowest weight or cost
- Dijkstra's Algorithm:
 - For a given source node in the graph, the algorithm finds the shortest path between that node and every other



Dijkstra's Algorithm


```
1: DIJKSTRA( $G, w, s$ )
2:    $d[s] \leftarrow 0$ 
3:   for all vertex  $v \in V$  do
4:      $d[v] \leftarrow \infty$ 
5:    $Q \leftarrow \{V\}$ 
6:   while  $Q \neq \emptyset$  do
7:      $u \leftarrow \text{EXTRACTMIN}(Q)$ 
8:     for all vertex  $v \in u.\text{ADJACENCYLIST}$  do
9:       if  $d[v] > d[u] + w(u, v)$  then
10:         $d[v] \leftarrow d[u] + w(u, v)$ 
```

Assignment Project Exam Help

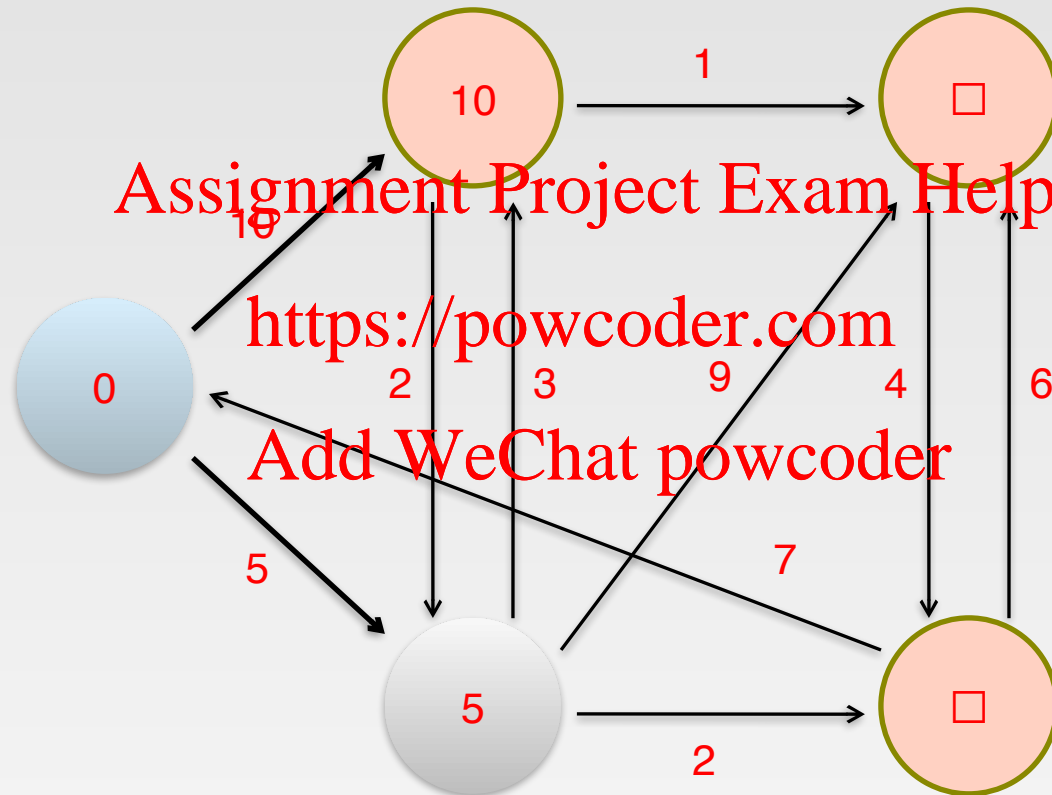
<https://powcoder.com>

Add WeChat powcoder

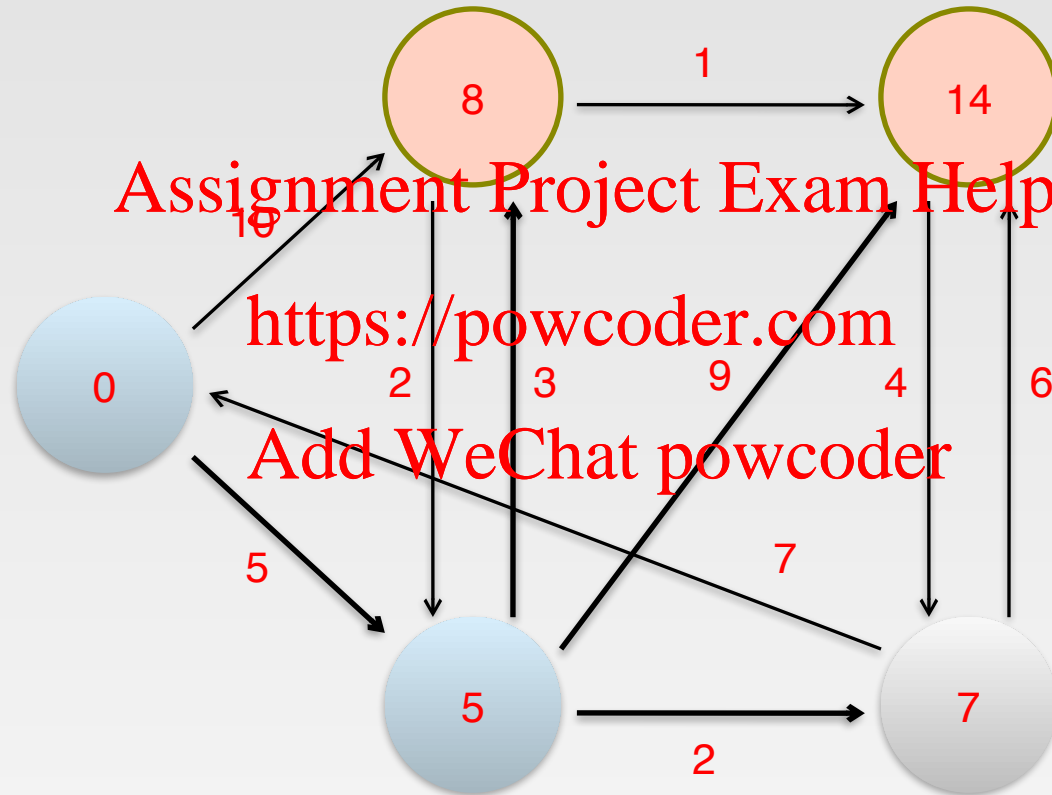
Dijkstra's Algorithm Example



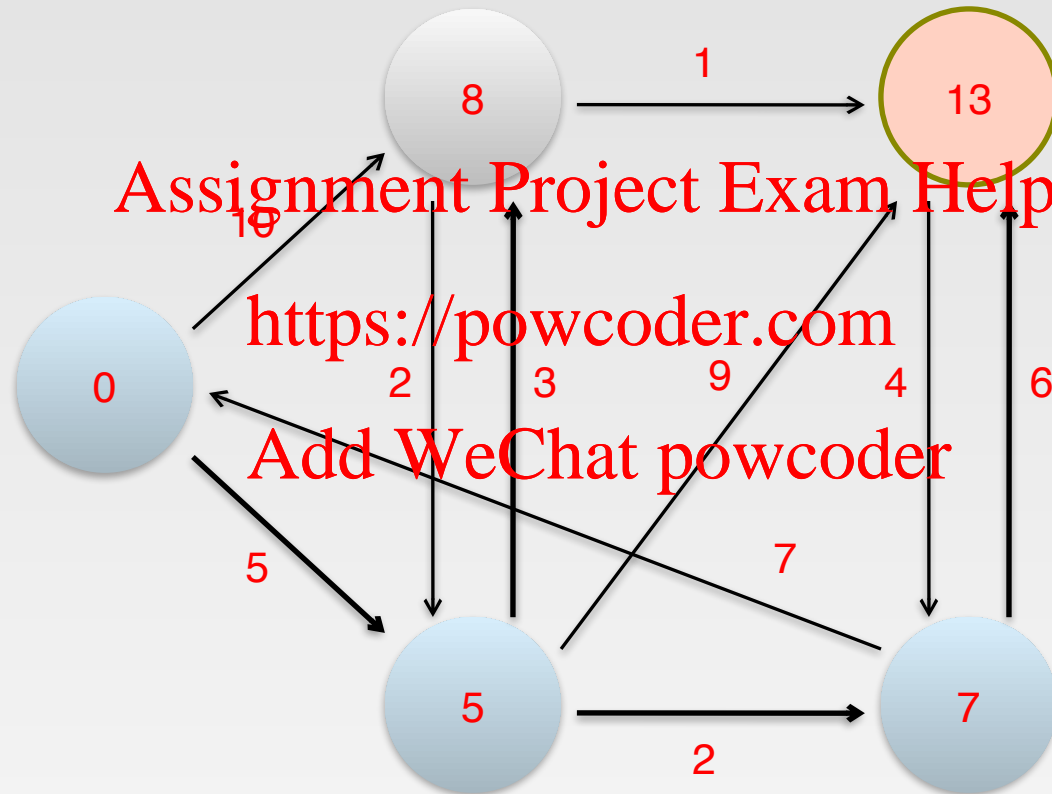
Dijkstra's Algorithm Example




Dijkstra's Algorithm Example




Dijkstra's Algorithm Example



Dijkstra's Algorithm Example



Dijkstra's Algorithm Example



Finish!

Single Source Shortest Path

- **Problem:** find shortest path from a source node to one or more target nodes
 - Shortest might also mean lowest weight or cost
- Single processor machine: Dijkstra's Algorithm
- MapReduce: parallel Breadth First Search (BFS)


Assignment Project Exam Help

<https://powcoder.com>


Add WeChat powcoder

Finding the Shortest Path

- Consider simple case of equal edge weights
- Solution to the problem can be defined inductively
- Here's the intuition:
 - Define: b is reachable from a if b is on adjacency list of a
 - $\text{DISTANCE}(s) = 0$
 - For all nodes p reachable from s , $\text{DISTANCE}(p) = 1$
 - For all nodes n reachable from some other set of nodes M , $\text{DISTANCE}(n) = 1 + \min(\text{DISTANCE}(m), m \in M)$



Visualizing Parallel BFS



From Intuition to Algorithm

- Data representation:
 - Key: node n
 - Value: d (distance from start), adjacency list (list of nodes reachable from n)
 - Initialization: for all nodes except for start node, $d = \infty$
- Mapper:
 - $\forall m \in \text{adjacency list: emit}(m, d + 1)$
- Sort/Shuffle
 - Groups distances by reachable nodes
- Reducer:
 - Selects minimum distance path for each reachable node
 - Additional bookkeeping needed to keep track of actual path

Multiple Iterations Needed

- Each MapReduce iteration advances the “known frontier” by one hop
 - Subsequent iterations include more and more reachable nodes as frontier expands
 - The input of Mapper is the output of Reducer in the previous iteration
 - Multiple iterations are needed to explore entire graph

Assignment Project Exam Help

<https://powcoder.com>

- Preserving graph structure:
 - Problem: Where did the adjacency list go?
 - Solution: mapper emits (n , adjacency list) as well

Add WeChat powcoder

BFS Pseudo-Code

- Equal Edge Weights (how to deal with weighted edges?)
- Only distances, no paths stored (how to obtain paths?)

```
class Mapper
  method Map(nid n, node N)
    d ← N.Distance
    Emit(nid n, N) //Pass along graph structure
    for all nodeid m ∈ N.AdjacencyList do
      Emit(nid m, d+1) //Emit distances to reachable nodes
```

```
class Reducer
  method Reduce(nid m, [d1, d2, ...])
    dmin ← ∞
    M ← ∅
    for all d ∈ counts [d1, d2, ...] do
      if IsNode(d) then
        M ← d //Recover graph structure
      else if d < dmin then //Look for shorter distance
        dmin ← d
    M.Distance ← dmin //Update shortest distance
    Emit(nid m, node M)
```

Stopping Criterion

- How many iterations are needed in parallel BFS (equal edge weight case)?
- Convince yourself: when a node is first “discovered”, we’ve found the shortest path
- Now answer the question.
 - The diameter of the graph, or the greatest distance between any pair of nodes
 - Six degrees of separation?
 - ▶ If this is indeed true, then parallel breadth-first search on the global social network would take at most six MapReduce iterations.

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Implementation in MapReduce

- The actual checking of the termination condition must occur outside of MapReduce.
- The driver (main) checks to see if a termination condition has been met, and if not, repeats.
- Hadoop provides a lightweight API called “counters”.
 - It can be used for counting events that occur during execution, e.g., number of corrupt records, number of times a certain condition is met, or anything that the programmer desires.
 - Counters can be designed to count the number of nodes that have distances of ∞ at the end of the job, the driver program can access the final counter value and check to see if another iteration is necessary.

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

How to Find the Shortest Path?

- The parallel breadth-first search algorithm only finds the shortest distances.
- Store “back-pointers” at each node, as with Dijkstra's algorithm
 - Not efficient to recover the path from the back pointers
- A simpler approach is to emit paths along with distances in the mapper, so that each node will have its shortest path easily accessible at all times
 - The additional space requirement is acceptable

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Mapper

```
public static class TheMapper extends Mapper<LongWritable, Text, LongWritable, Text> {  
    public void map(LongWritable key, Text value, Context context) throws IOException, InterruptedException {  
        Text word = new Text();  
        String line = value.toString();//looks like 1 0 2:3:  
        String[] sp = line.split(" ");//splits on space  
        int distanceadd = Integer.parseInt(sp[1]) + 1;  
        String[] PointsTo = sp[2].split(":");  
        for(int i=0; i<PointsTo.length; i++){  
            word.set("VALUE "+distanceadd);//tells me to look at distance value  
            context.write(new LongWritable(Integer.parseInt(PointsTo[i])), word);  
            word.clear(); }  
        //pass in current node's distance (if it is the lowest distance)  
        word.set("VALUE "+sp[1]);  
        context.write( new LongWritable( Integer.parseInt( sp[0] ) ), word );  
        word.clear();  
        word.set("NODES "+sp[2]);//tells me to append on the final tally  
        context.write( new LongWritable( Integer.parseInt( sp[0] ) ), word );  
        word.clear();  
    }  
}
```

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Reducer

```
public static class TheReducer extends Reducer<LongWritable, Text, LongWritable, Text> {  
    public void reduce(LongWritable key, Iterable<Text> values, Context context) throws  
IOException, InterruptedException {  
        String nodes = "UNMODED";  
        Text word = new Text();  
        int lowest = INFINITY; //start at infinity
```

```
        for (Text val : values) {  
            //looks like NODES/VALUES 1 0 2:3:, we need to use the first as a key
```

```
            String[] sp = val.toString().split(" "); //splits on space
```

```
            //look at first value https://powcoder.com
```

```
            if(sp[0].equalsIgnoreCase("NODES")){
```

```
                nodes = null;
```

```
                nodes = sp[1]; Add WeChat powcoder
```

```
            }else if(sp[0].equalsIgnoreCase("VALUE")){
```

```
                int distance = Integer.parseInt(sp[1]);
```

```
                lowest = Math.min\(distance, lowest\);
```

```
            }
```

```
        }
```

```
        word.set(lowest+" "+nodes);
```

```
        context.write(key, word);
```

```
        word.clear();
```

```
    }
```

```
}
```

<https://github.com/himank/Graph-Algorithm-Map-Reduce/blob/master/src/DijkstraAlgo.java>

BFS Pseudo-Code (Weighted Edges)

- The adjacency lists, which were previously lists of node ids, must now encode the edge distances as well
 - Positive weights!
- In line 6 of the mapper code, instead of emitting $d+1$ as the value, we must now emit $d + w$, where w is the edge distance

<https://powcoder.com>


- **The termination behaviour is very different!**

- How many iterations are needed in parallel BFS (positive edge weight case)?

- Convince yourself: when a node is first “discovered”, we’ve found the shortest path

Not true!

Additional Complexities




Assignment Project Exam Help

- Assume that p is the current processed node
 - In the current iteration, we just “discovered” node r for the very first time.
 - We've already discovered the shortest distance to node p , and that the shortest distance to r so far goes through p
 - Is $s \rightarrow p \rightarrow r$ the shortest path from s to r ?
- The shortest path from source s to node r may go outside the current search frontier
 - It is possible that $p \rightarrow q \rightarrow r$ is shorter than $p \rightarrow r$!
 - We will not find the shortest distance to r until the search frontier expands to cover q .

How Many Iterations Are Needed?

- In the worst case, we might need as many iterations as there are nodes in the graph minus one
 - A sample graph that elicits worst-case behaviour for parallel breadth-first search.
 - Eight iterations are required to discover shortest distances to all nodes from n_1 .



Example (only distances)

□ Input file:

s --> 0 | n1: 10, n2: 5

n1 --> ∞ | n2: 2, n3:1

n2 --> ∞ | n1: 3, n3:9 , n4:2

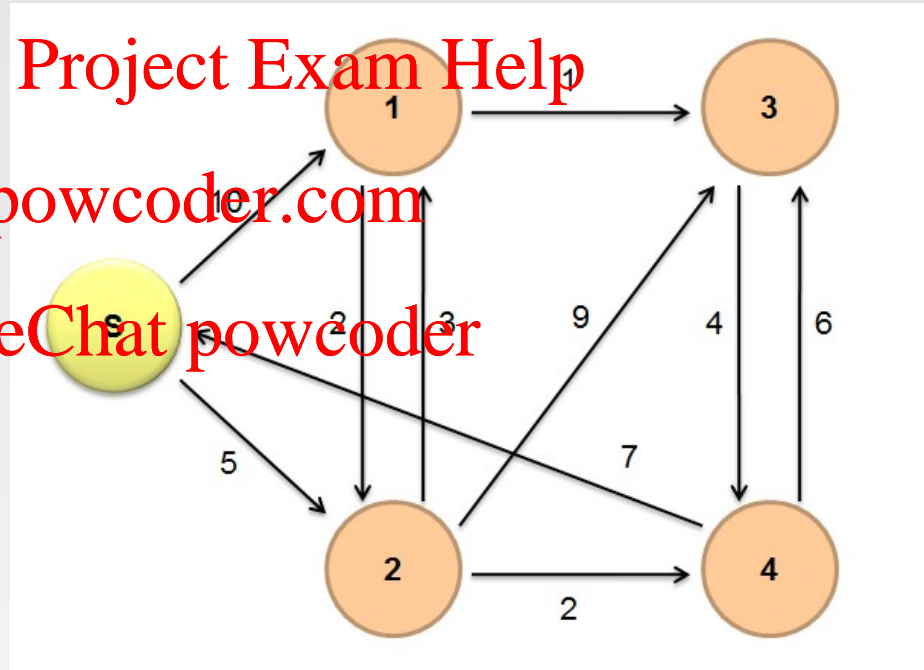
n3 --> ∞ | n4:4

n4 --> ∞ | s:7, n3:6

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



Iteration 1

□ Map:

Read s --> 0 | n1: 10, n2: 5

Emit: (n1, 10), (n2, 5), and the adjacency list (s, n1: 10, n2: 5)

The other lists will also be read and emit, but they do not contribute, and thus ignored

Assignment Project Exam Help

□ Reduce:

Receives: (n1, 10), (n2, 5), (s, 10), (s, 5)


The adjacency list of each node will also be received, ignored in example

Emit: Add WeChat powcoder

s --> 0 | n1: 10, n2: 5

n1 --> 10 | n2: 2, n3:1

n2 --> 5 | n1: 3, n3:9 , n4:2



Iteration 2

□ Map:

Read: n1 --> 10 | n2: 2, n3:1

Emit: (n2, 12), (n3, 11), (n1, <10, (n2: 2, n3:1)>)

Read: n2 --> 5 | n1: 3, n3:9 , n4:2

Emit: (n1, 8), (n3, 14), (n4, 7), (n2, <5, (n1:3, n3:9, n4:2)>)

Ignore the processing of the other lists

□ Reduce:

Receives: (n1, (8, <10, (n2: 2, n3:1)>)), (n2, (12, <5, n1: 3, n3:9 , n4:2>)), (n3, (11, 14)), (n4, 7)


Emit:

n1 --> 8 | n2: 2, n3:1

n2 --> 5 | n1: 3, n3:9 , n4:2

n3 --> 11 | n4:4

n4 --> 7 | s:7, n3:6



Iteration 3

□ Map:

Read: n1 --> 8 | n2: 2, n3:1

Emit: (n2, 10), (n3, 9), (n1, <8, (n2: 2, n3:1)>)

Read: n2 --> 5 | n1: 3, n3:9 , n4:2 (**Again!**)

Emit: (n1, 8), (n3, 14), (n4, 7), (n2, 5), (n1:3, n3:9, n4:2)>)

Read: n3 --> 11 | n4:4

Emit: (n4, 15) , (n3, <11, (n4:4)>)

Read: n4 --> 7 | s:7, n3:6

Emit: (s, 14), (n3, 13), (n4, <7, (s:7, n3:6)>)

□ Reduce:


Emit:

n1 --> 8 | n2: 2, n3:1

n2 --> 5 | n1: 3, n3:9 , n4:2

n3 --> 9 | n4:4

n4 --> 7 | s:7, n3:6



Iteration 4

□ Map:

Read: n1 --> 8 | n2: 2, n3:1 (**Again!**)

Emit: (n2, 10), (n3, 9), (n1, <8, (n2: 2, n3:1)>)

Read: n2 --> 5 | n1: 3, n3:9 , n4:2 (**Again!**)

Emit: (n1, 8), (n3, 14), (n4, 7), (n2, <5, (n1:3, n3:9, n4:2)>)

Read: n3 --> 9 | n4:4

Emit: (n4, 13) , (n3, <9, (n4:4)>)

Read: n4 --> 7 | s:7, n3:6 (**Again!**)

Emit: (s, 14), (n3, 13), (n4, <7, (s:7, n3:6)>)

□ Reduce:

Emit:


n1 --> 8 | n2: 2, n3:1

n2 --> 5 | n1: 3, n3:9 , n4:2

n3 --> 9 | n4:4

n4 --> 7 | s:7, n3:6

In order to avoid duplicated computations, you can use a status value to indicate whether the distance of the node has been modified in the previous iteration.



No updates. Terminate.

Comparison to Dijkstra

- Dijkstra's algorithm is more efficient
 - At any step it only pursues edges from the minimum-cost path inside the frontier
- MapReduce explores all paths in parallel
 - Lots of “waste”
 - Useful work is only done at the “frontier”
- Why can't we do better using Map Reduce?

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Graphs and MapReduce

- Graph algorithms typically involve:
 - Performing computations at each node: based on node features, edge features, and local link structure
 - Propagating computations: “traversing” the graph
- Generic recipe:
 - Represent graphs as adjacency lists
 - Perform local computations in mapper
 - Pass along partial results via outlinks, keyed by destination node
 - Perform aggregation in reducer on inlinks to a node
 - Iterate until convergence: controlled by external “driver”
 - Don’t forget to pass the graph structure between iterations

Assignment Project Exam Help

<https://powcoder.com>


Add WeChat powcoder

Issues with MapReduce on Graph Processing

- MapReduce Does not support iterative graph computations:
 - External driver. Huge I/O incurs
 - No mechanism to support global data structures that can be accessed and updated by all mappers and reducers
 - ▶ Passing information is only possible within the local graph structure – through adjacency list
 - ▶ Dijkstra's algorithm on a single machine: a global priority queue that guides the expansion of nodes
 - ▶ Dijkstra's algorithm in Hadoop: no such queue available. Do some “wasted” computation instead
- MapReduce algorithms are often impractical on large, dense graphs.
 - The amount of intermediate data generated is on the order of the number of edges.
 - For dense graphs, MapReduce running time would be dominated by copying intermediate data across the network.

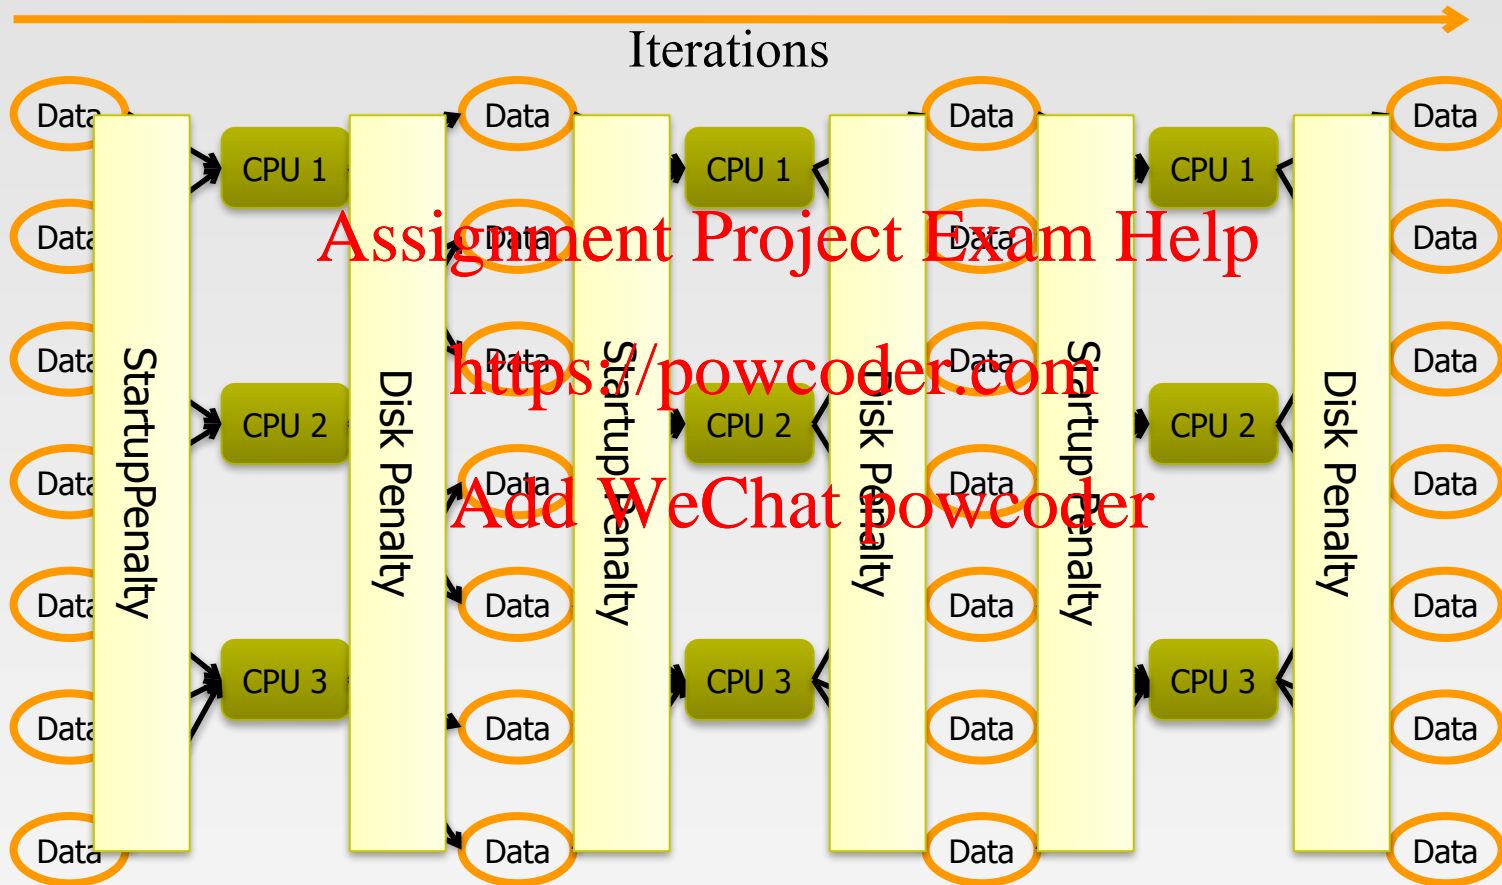
Iterative MapReduce

- Only a subset of data needs computation:



Iterative MapReduce

- System is not optimized for iteration:



Better Partitioning

- Default: hash partitioning
 - Randomly assign nodes to partitions
- Observation: many graphs exhibit local structure
 - E.g., communities in social networks
 - Better partitioning creates more opportunities for local aggregation
- Unfortunately, partitioning is **hard**!
 - Sometimes, chick-and-egg...
 - But cheap heuristics sometimes available
 - For webgraphs: range partition on domain-sorted URLs

Assignment Project Exam Help

<https://powcoder.com>


Add WeChat powcoder

Schimmy Design Pattern

- Basic implementation contains two dataflows:
 - Messages (actual computations)
 - Graph structure (“bookkeeping”)
- Schimmy: separate the two dataflows, shuffle only the messages
 - Basic idea: merge join between graph structure and messages

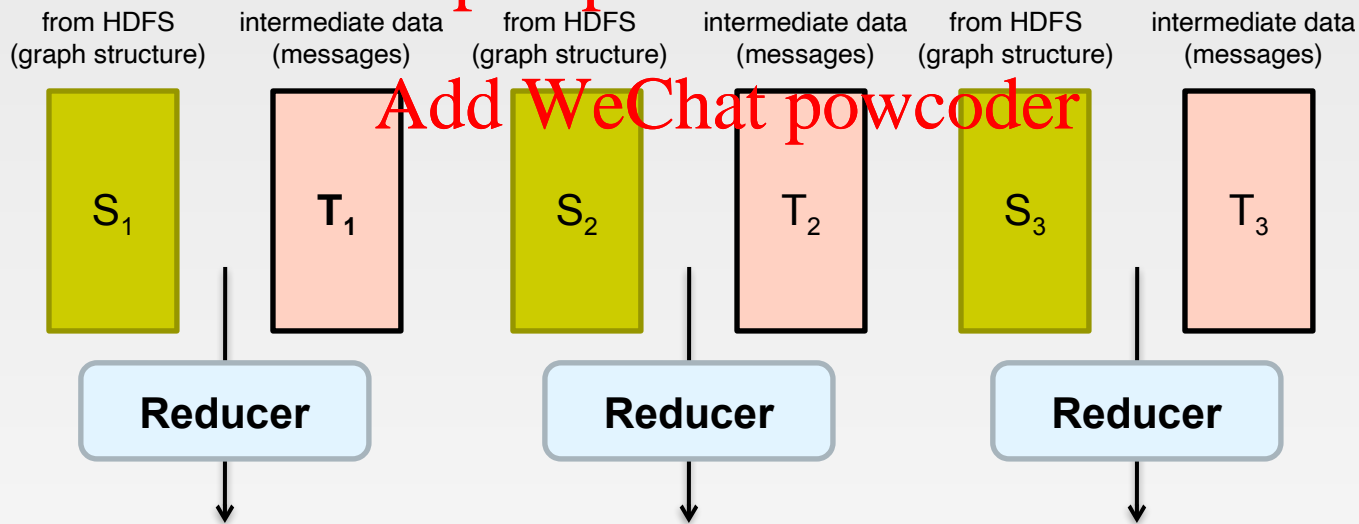
<https://powcoder.com>

both relations are consistently partitioned and sorted by join key



Do the Schimmy!

- Schimmy = reduce side parallel merge join between graph structure and messages
 - Consistent partitioning between input and intermediate data
 - Mappers emit only messages (actual computation)
 - Reducers read graph structure directly from HDFS



Assignment Project Exam Help

<https://powcoder.com>
Introduction to Pregel

Add WeChat powcoder

Motivation of Pregel

- Many practical computing problems concern large graphs

Large graph data

Web graph
Transportation routes
Citation relationships
Social networks

Graph algorithms

PageRank
Shortest path
Connected components
Clustering techniques



Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

- Single computer graph library does not scale
- MapReduce is ill-suited for graph processing
 - Many iterations are needed for parallel graph processing
 - Materializations of intermediate results at every MapReduce iteration harm performance

Pregel

- **Pregel**: A System for Large-Scale **Graph** Processing (Google) - Malewicz et al. SIGMOD 2010.

- Scalable and Fault-tolerant platform

Assignment Project Exam Help

- API with flexibility to express arbitrary algorithm

<https://powcoder.com>

- Inspired by Valiant's Bulk Synchronous Parallel model

- Leslie G. Valiant: A Bridging Model for Parallel Computation. Commun. ACM 33 (8): 103-111 (1990)

- Vertex centric computation (Think like a vertex)

Bulk Synchronous Parallel Model (BSP)

analogous to MapReduce rounds

- Processing: a series of **supersteps**
- **Vertex**: computation is defined to run on each vertex
- **Superstep S**: *all vertices compute in parallel; each vertex v may*
 - receive **messages** sent to v from superstep S – 1;
 - perform some computation: modify its states and the states of its outgoing edges
 - Send **messages** to other vertices (to be received in the next superstep)

Add WeChat powcoder

Vertex-centric, message passing

Pregel Computation Model


- Based on Bulk Synchronous Parallel (BSP)
 - Computational units encoded in a directed graph
 - Computation proceeds in a series of supersteps
 - Message passing architecture

Input



<https://powcoder.com>

Add WeChat powcoder



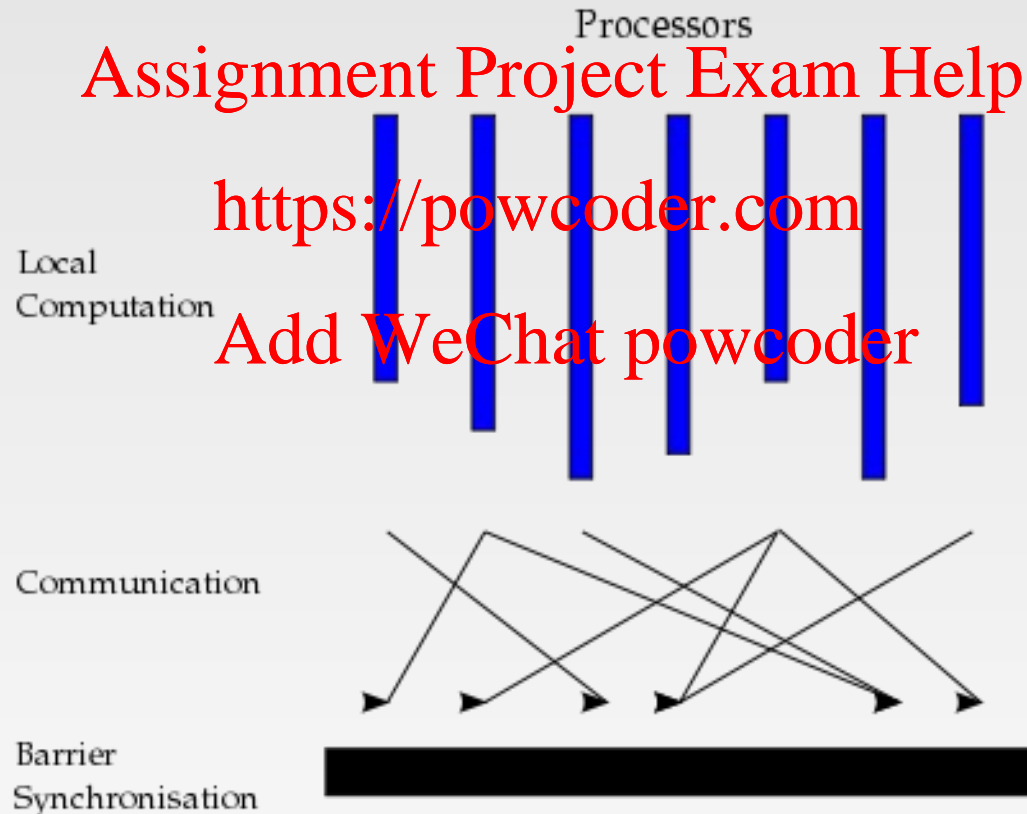
Supersteps
(a sequence of iterations)



Output

Pregel Computation Model (Cont')

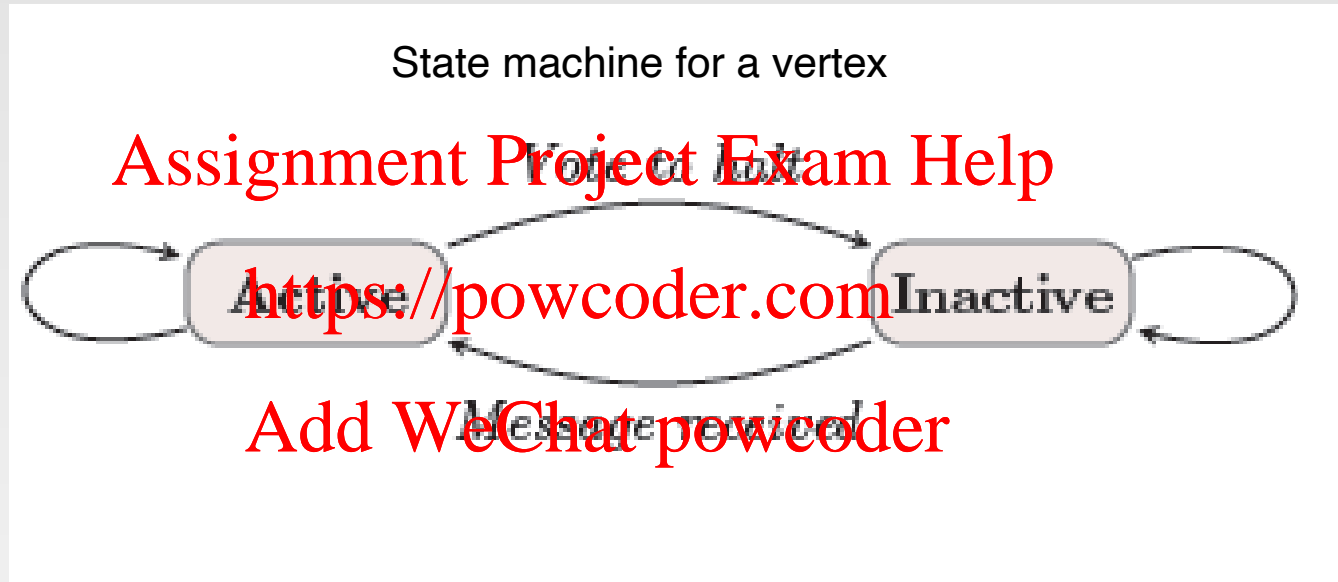
- Concurrent computation and Communication need not be ordered in time
- Communication through message passing



Source: http://en.wikipedia.org/wiki/Bulk_synchronous_parallel


Pregel Computation Model (Cont')

- Superstep: the vertices compute in parallel
- Each vertex



- Termination condition
 - ▶ All vertices are simultaneously inactive
 - ▶ A vertex can choose to deactivate itself
 - ▶ Is “woken up” if new messages received

Example: SSSP – Parallel BFS in Pregel



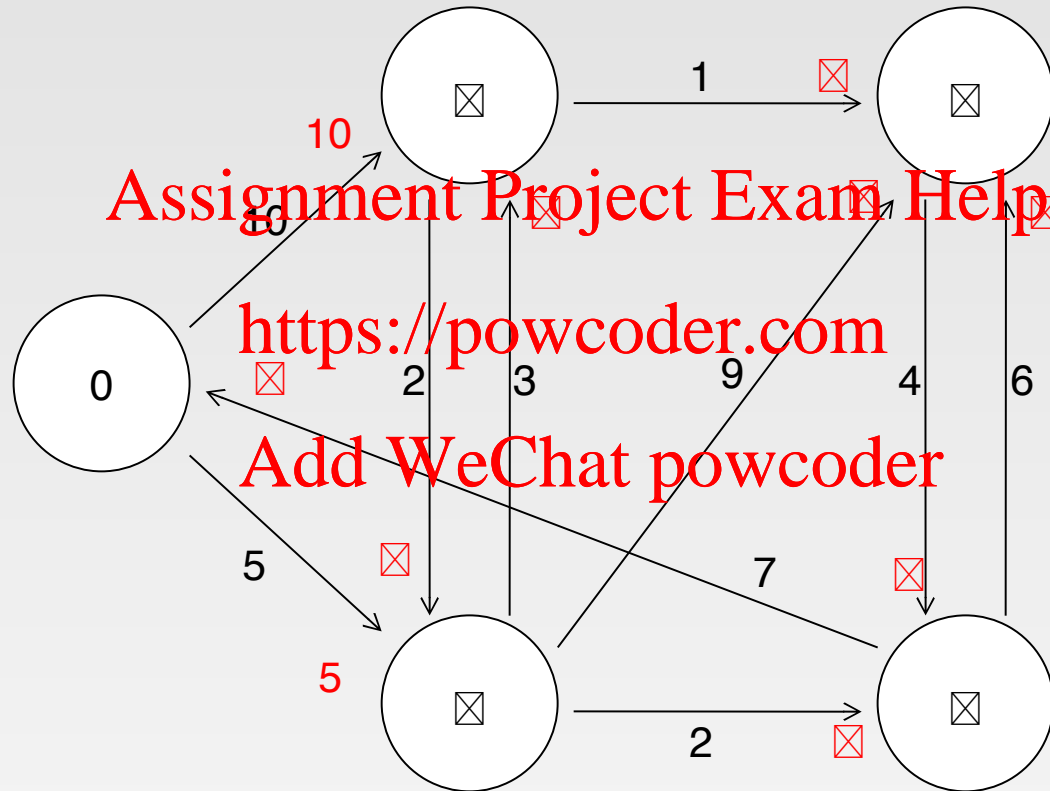
● Inactive Vertex

○ Active Vertex

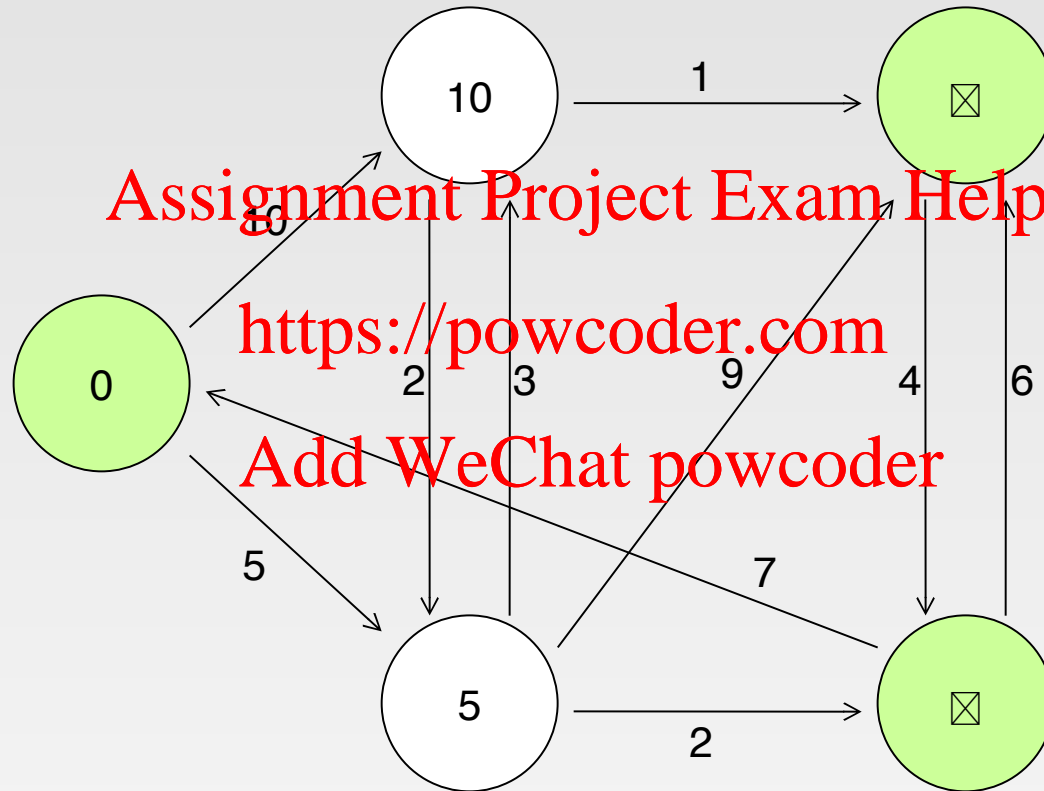
\xrightarrow{x} Edge weight

\times Message

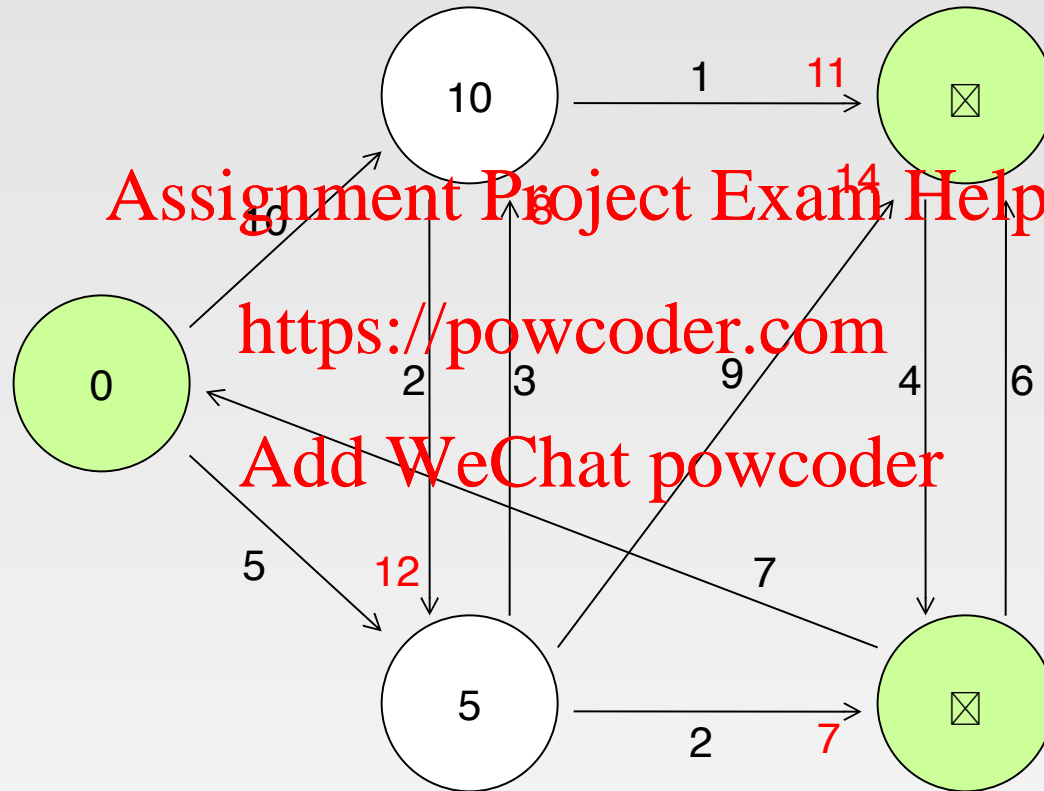
Example: SSSP – Parallel BFS in Pregel



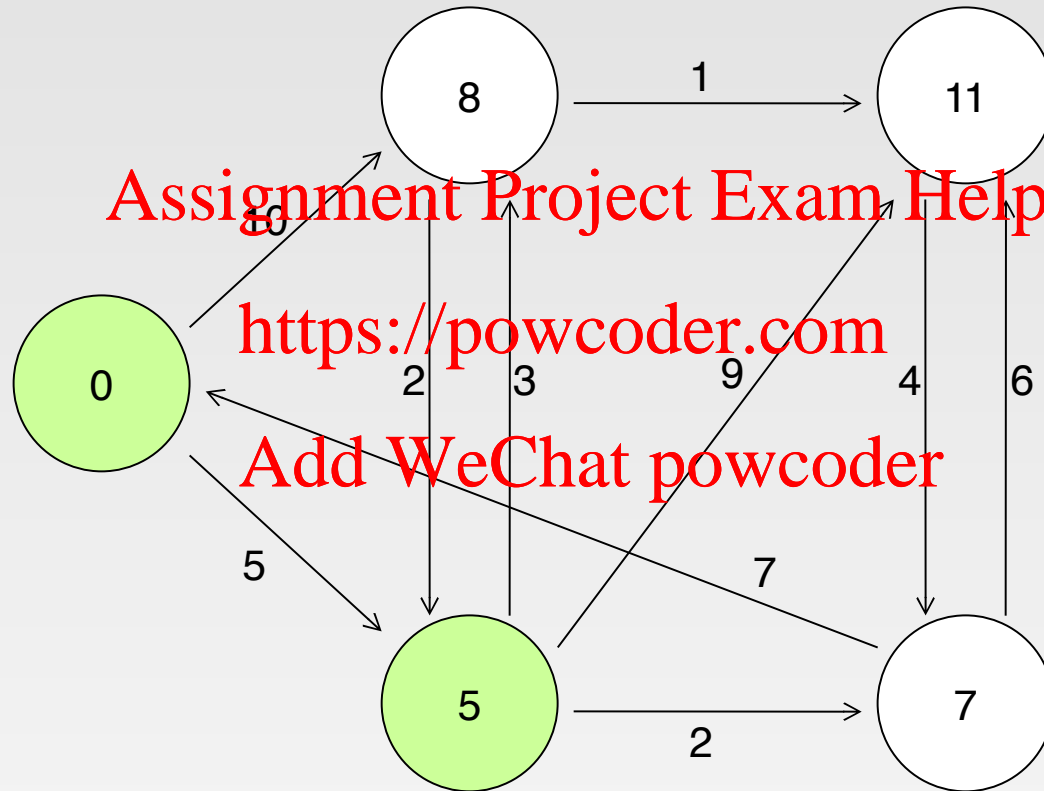
Example: SSSP – Parallel BFS in Pregel



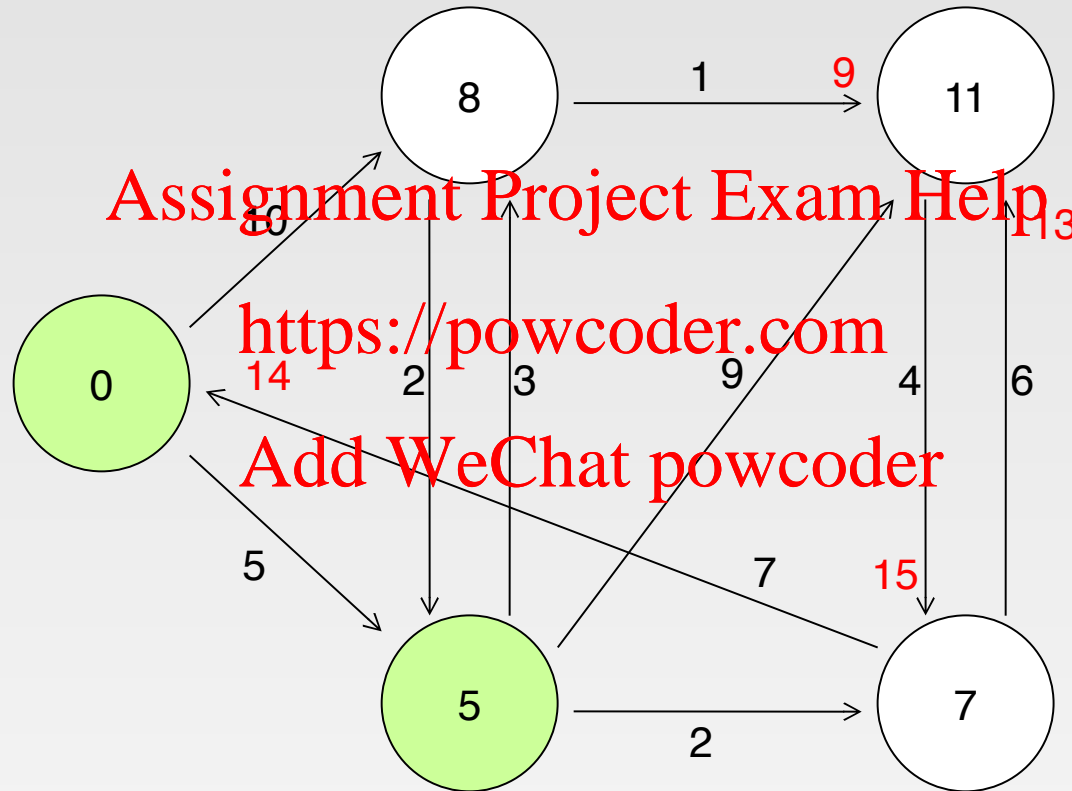
Example: SSSP – Parallel BFS in Pregel



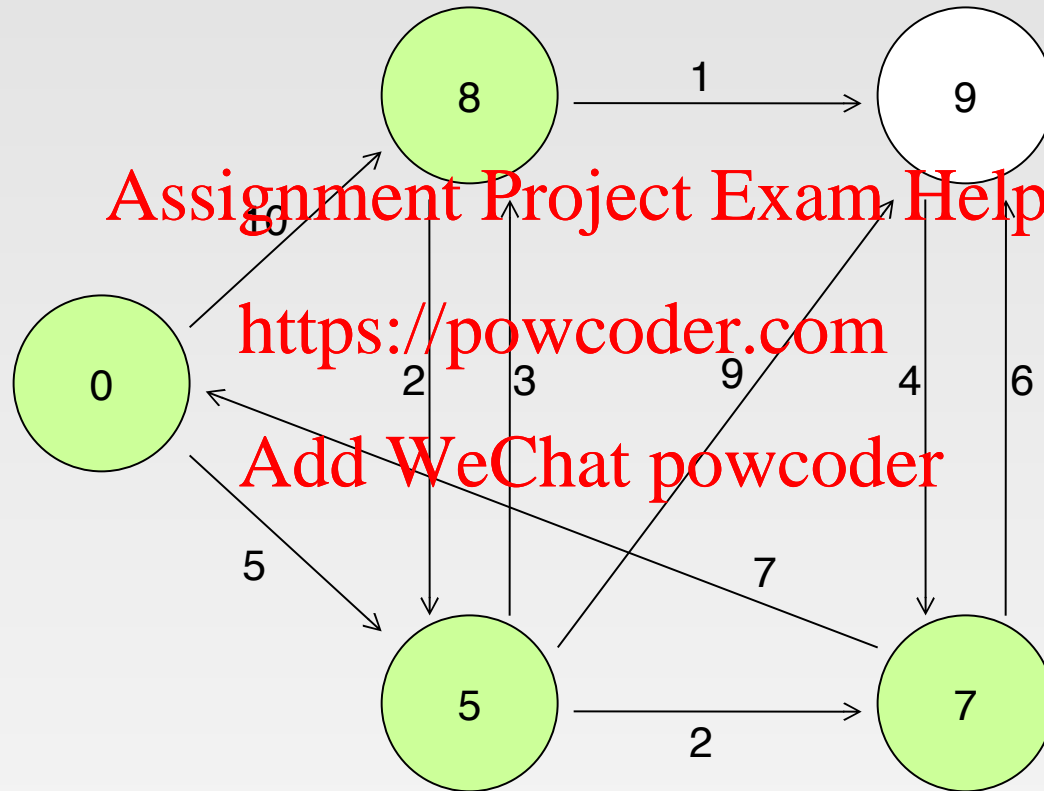
Example: SSSP – Parallel BFS in Pregel



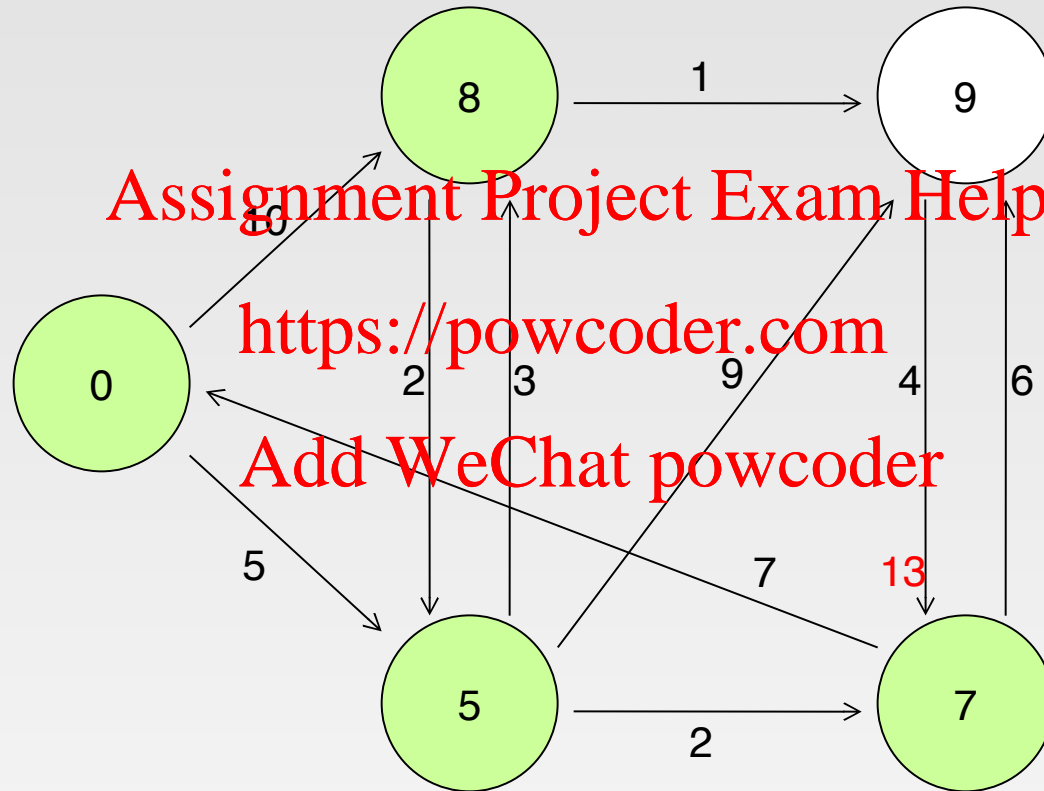
Example: SSSP – Parallel BFS in Pregel



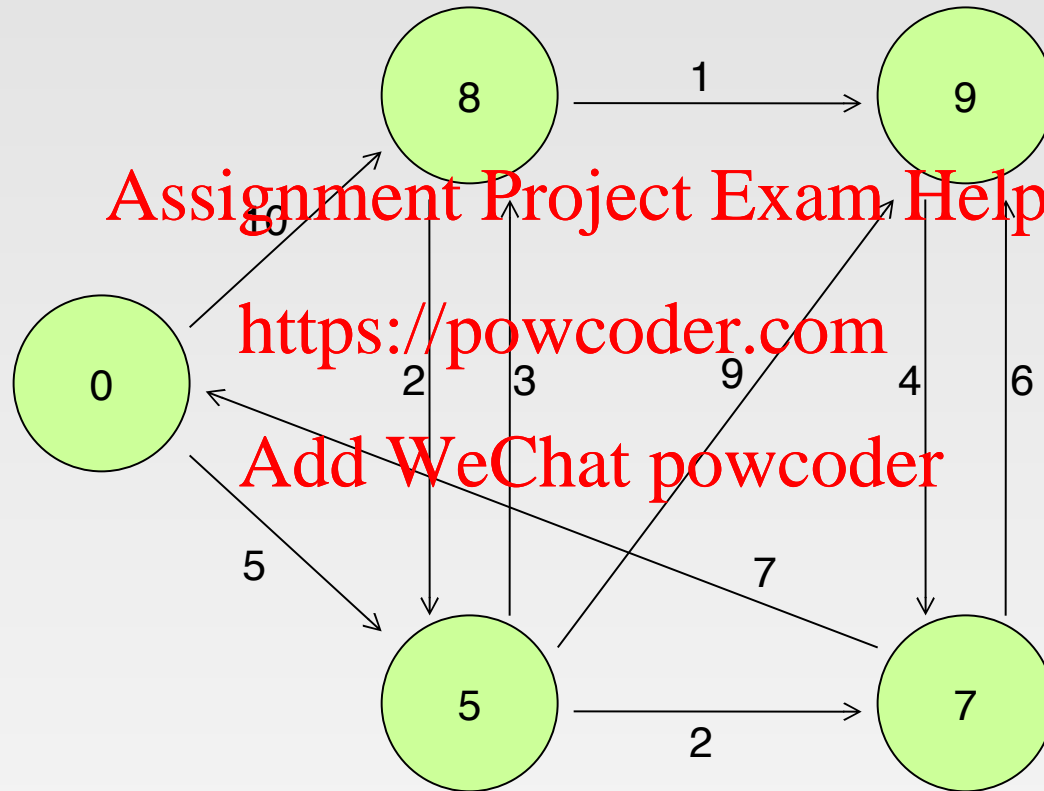
Example: SSSP – Parallel BFS in Pregel



Example: SSSP – Parallel BFS in Pregel



Example: SSSP – Parallel BFS in Pregel



Differences from MapReduce

- Graph algorithms can be written as a series of chained MapReduce jobs
- Pregel
 - Keeps vertices & edges on the machine that performs computation
 - Uses network transfers only for messages
- MapReduce
 - Passes the entire state of the graph from one stage to the next
 - Needs to coordinate the steps of a chained MapReduce

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Writing a Pregel Program (C++)

- Subclassing the predefined **Vertex** class

```
template <typename VertexValue,
          typename EdgeValue,
          typename MessageValue>
class Vertex {
public:
    virtual void Compute(MessageIterator* msgs) = 0;
    const string& vertex_id() const;
    int64 superstep() const;
    const VertexValue& GetValue();
    VertexValue* MutableValue();
    OutEdgeIterator GetOutEdgeIterator();
    void SendMessageTo(const string& dest_vertex,
                      const MessageValue& message);
    void VoteToHalt();
};
```

Override this!

in msgs

Modify vertex value

out msg

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Pregel: SSSP (C++)

Refer to the current node as u

```
class ShortestPathVertex : public Vertex<int, int, int>
{
    void Compute(MessageIterator* msgs) {
        int mindist = IsSource(vertex_id()) ? 0 : INF;
        for (; !msgs->Done(); msgs->Next())
            mindist = min(mindist, msgs->MutableValue());
        if (mindist < GetValue())
            *MutableValue() = mindist;
        OutEdgeIterator iter = GetOutEdgeIterator();
        for (; !iter.Done(); iter.Next())
            SendMessageTo(iter.Target(), mindist + iter.GetValue());
    }
    VoteToHalt();
};
```

aggregation

Assignment Project Exam Help

Messages: distances to u

<https://powcoder.com>

MutableValue: the current distance

Add WeChat powcoder

Pass revised distance to its neighbors

More Tools on Big Graph Processing

□ Graph databases: Storage and Basic Operators

- http://en.wikipedia.org/wiki/Graph_database
- Neo4j (an open source graph database)
- InfiniteGraph
- VertexDB
-

Assignment Project Exam Help

<https://powcoder.com>

□ Distributed Graph Processing (mostly in-memory-only)

- Google's Pregel (vertex centered computation)
- Giraph (Apache)
- GraphX (Spark)
- GraphLab
-

Add WeChat powcoder

References

- Chapter 5. Data-Intensive Text Processing with MapReduce
- Chapter 5. Mining of Massive Datasets.

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Assignment Project Exam Help

End of Chapter 5


<https://powcoder.com>

Add WeChat powcoder

PageRank Review

- Given page t_j with in-coming neighbors $t_1 \dots t_n$, where
 - d_i is the out-degree of t_i
 - β is the teleport probability
 - N is the total number of nodes in the graph

Assignment Project Exam Help



Computing PageRank

- Properties of PageRank
 - Can be computed iteratively
 - Effects at each iteration are local
- Sketch of algorithm:
 - Start with seed r_i values
 - Each page distributes r_i “credit” to all pages it links to
 - Each target page t_j adds up “credit” from multiple in-bound links to compute r_j
 - Iterate until values converge

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Simplified PageRank

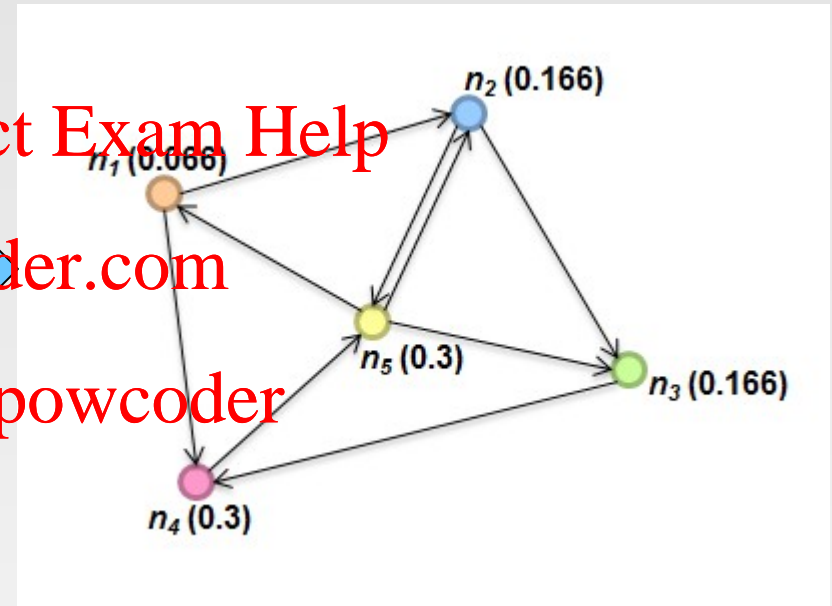
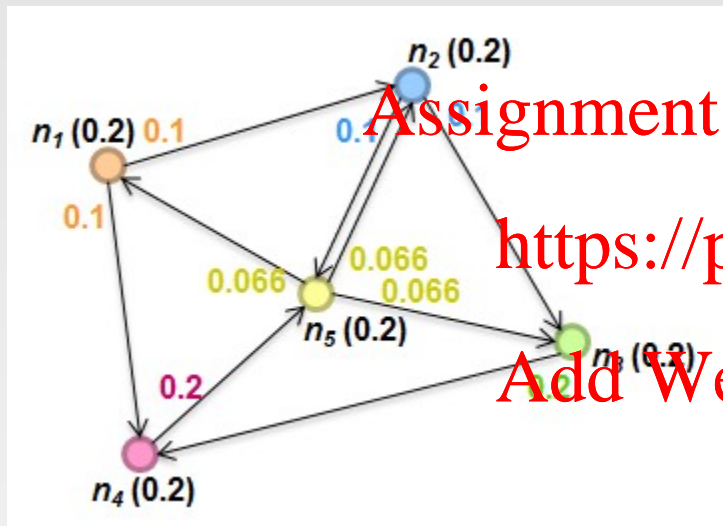
- First, tackle the simple case:
 - No teleport
 - No dangling nodes (dead ends)
 - Then, factor in these complexities...
- How to deal with the teleport probability?
 - How to deal with dangling nodes?

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Sample PageRank Iteration (1)





Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Sample PageRank Iteration (2)


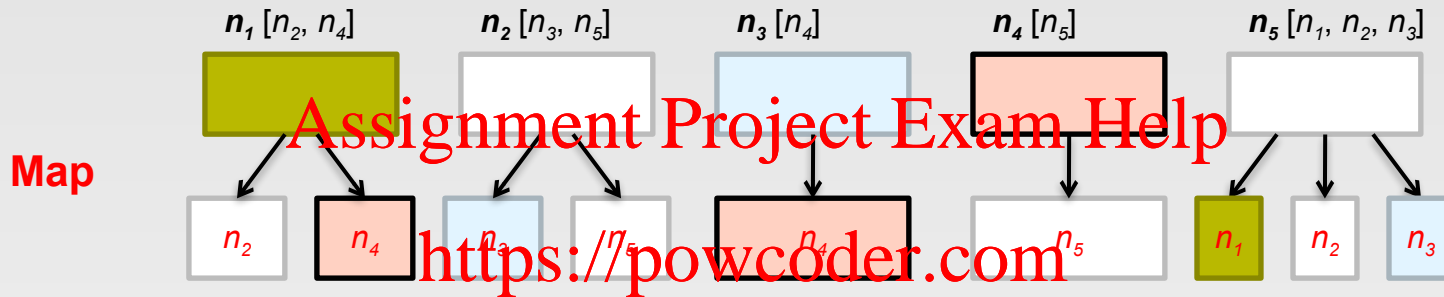


Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

PageRank in MapReduce (One Iteration)



PageRank Pseudo-Code

```
1: class MAPPER
2:   method MAP(nid n, node N)
3:      $p \leftarrow N.PAGERANK / |N.ADJACENCYLIST|$ 
4:     EMIT(nid n, N) ▷ Pass along graph structure
5:     for all nodes  $m \in N.ADJACENCYLIST$  do
6:       EMIT(nid m, p) ▷ Pass PageRank mass to neighbors

1: class REDUCER
2:   method REDUCE(nid m, [ $p_1, p_2, \dots$ ])
3:      $M \leftarrow \emptyset$ 
4:     for all  $p \in counts[p_1, p_2, \dots]$  do
5:       if ISNODE(p) then
6:          $M \leftarrow p$  ▷ Recover graph structure
7:       else
8:          $s \leftarrow s + p$  ▷ Sums incoming PageRank contributions
9:      $M.PAGERANK \leftarrow s$ 
10:    EMIT(nid m, node M)
```

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

PageRank vs. BFS

	PageRank	BFS
Assignment Project Exam Help		
Map	r_i/d_i	$d+w$
Reduce	sum	min

<https://powcoder.com>
Add WeChat powcoder

PageRank in Pregel

- **Superstep 0:** Value of each vertex is `1/NumVertices()`

```
virtual void Compute(MessageIterator* msgs) {  
    if (superstep() >= 1) {  
        double sum = 0;  
        for (; !msgs->done(); msgs->Next())  
            sum += msgs->Value();  
        *MutableValue() = 0.15 + 0.85 * sum;  
    }  
    if (supersteps() < 30) {  
        const int64 n = GetOutEdgeIterator().size();  
        SendMessageToAllNeighbors(GetValue() / n);  
    } else {  
        VoteToHalt();  
    }  
}
```